

Distributed Raft Cache

Consensus-Driven Caching with ML-Based Eviction

Problem & Solution

The Problem

- ✗ Single-node caches fail completely, leading to **service outages**
- ✗ Network partitions cause **data inconsistency** across replicas
- ✗ LRU eviction **doesn't adapt** to complex access patterns
- ✗ Difficult trade-off: **Speed vs Consistency vs Fault-tolerance**

Our Solution

01



Raft Consensus

Strong consistency guarantees across distributed nodes

02



Read Lease

105x faster reads by bypassing log replication

03



ML Eviction

Smart cache eviction using Random Forest prediction


04




Demo UI

Real-time visualization of cluster state & elections

System Architecture



Tech Stack



Raft Consensus Algorithm

⌚ Primary Goals

- Enable multiple servers to agree on shared state reliably
- Ensure system functionality even if minority of servers fail
- Guarantee strong data consistency across the cluster

蓄 Three Sub-problems

1

Leader Election

Selecting a coordinator to manage the log replication

2

Log Replication

Accepting entries and replicating them across the cluster

3

Safety

Ensuring state machines apply entries in the same order

Three States of a Raft Node



STATE	ROLE
● FOLLOWER	Passive - responds to requests from leaders and candidates
● CANDIDATE	Active - seeking votes from other nodes to become leader
● LEADER	Dominant - handles all client read/write requests & replicates log


Leader Election Algorithm

● ● ● raft_election.log

LEADER ELECTION PROTOCOL

```
01 // 1. Initial State
02 Follower waits for heartbeat from Leader
03 // 2. On Timeout (150-300ms randomized)
04 if (timeout_elapsed) {
05   state = CANDIDATE;
06   term++;
07   voteFor(SELF);
08   requestVotes(ALL_NODES);
09 }
10 // 3. Election Result
11 if (votes > majority) {
12   state = LEADER;
13 } else if (higher_term_seen) {
14   state = FOLLOWER;
15 }
```

RANDOMIZED TIMEOUTS



Different timeout durations ensure one node triggers election before others, reducing collision probability.



KEY INSIGHT

"Randomized timeout prevents split votes"

Log Replication Algorithm

REPLICATION_PROTOCOL

</>

1. Client sends `write request` to Leader
2. Leader appends entry to local log
3. Leader sends `AppendEntries` RPC to followers
4. Wait for `majority` confirmation
5. Mark entry as **COMMITTED**
6. Apply command to state machine
7. Respond success to client

LEADER LOG STATE

INDEX	TERM	COMMAND	STATUS
1	T1	x ← 1	✓ Committed
2	T1	y ← 2	✓ Committed
3	T2	z ← 3	✓ Committed
4	T2	w ← 4	○ Replicating...

commitIndex = 3

ⓘ Entry 4 is currently being replicated to followers. It will be marked committed once a majority acknowledge receipt.

Raft Safety Guarantees

Five fundamental properties ensuring system correctness



Election Safety

At most one leader can be elected in a given term.



Leader Completeness

If log entry is committed in a term, it will be present in logs of leaders for all higher terms.



Leader Append-Only

A leader never overwrites or deletes entries in its log; it only appends new entries.



State Machine Safety

If a server has applied a log entry at a given index, no other server will ever apply a different command for the same index.



Log Matching

If two logs contain an entry with the same index and term, then the logs are identical in all entries up through the given index.

THE GUARANTEE

"All nodes eventually have identical data"

Handling Failure



Follower Recovery

⚡ Leader Failure Timeline

T=0.00s

Leader Node crashes ✗

Process termination or network partition

T=0.20s

Followers detect timeout

No heartbeats received for 200ms

T=0.35s

Node2 becomes CANDIDATE

Increments term, votes for self

T=0.40s

Node2 gets majority votes

Received vote from Node3

T=0.41s

Node2 is new LEADER ✓

Cluster operational again

TOTAL DOWNTIME

~500ms

1 Follower reconnects

Node rejoins the network after restart

2

Leader detects lag

Compares nextIndex with follower log

3

Send missing entries

Leader streams AppendEntries RPCs

4

Follower catches up

Applies all committed entries to state

5

Consistency Restored

Node is fully synchronized and serving

TYPICAL RECOVERY TIME


~10 seconds

Why Reads Are Slow

NAIVE APPROACH

"Route ALL reads through Raft log"

- ⚠ Every read requires **log append + replication** round-trip
- ⌚ Read latency degrades to Write latency (~50-100ms vs optimal <1ms)
- ⚡ Massive inefficiency: Reads are **80%** of the total workload!



OUR GOAL

"Fast reads WITHOUT sacrificing consistency"

ReadIndex Protocol

● ○ ● raft_read.log

READINDEX PROTOCOL

```
01 // 1. Capture State
02 readIndex = commitIndex;
03 // 2. Verify Leadership
04 broadcastHeartbeat();
05 wait(majority_response); // Confirms still leader
06 // 3. Wait for State Machine
07 while (lastApplied < readIndex) {
08   wait(); // Ensure data is up-to-date
09 }
10 // 4. Local Read & Return
11 result = stateMachine.lookup(key);
12 return result;
```

PROTOCOL BENEFITS



No Log Append Needed

Bypasses disk I/O for read operations



Strong Consistency

Guarantees no stale reads are returned



Moderate Latency

~50-100ms (1 RTT required)



KEY INSIGHT

"ReadIndex avoids the heavy Raft log replication path while preserving linearizability."

Read Lease - 105x Faster!

read_lease.py

READ LEASE OPTIMIZATION

```
01 // 1. Configuration
02 lease_duration = election_timeout / 2; // 75ms
03 // 2. Extend Lease on Heartbeat
04 if (heartbeat_ack_majority) {
05     lease_expiry = now() + 75ms;
06 }
07 // 3. Handle Read Request
08 function handleRead(key) {
09     if (now() < lease_expiry) {
10         return readLocally(key); // FAST!
11     } else {
12         return useReadIndex(key); // SAFE
13     }
14 }
```



KEY INSIGHT

"If we recently confirmed leadership, we are guaranteed to still be the leader."



WHY IT'S SAFE

Lease (75ms) < Election Timeout (150ms).
No other node can become leader while our lease is valid.



**0.5ms vs 52ms = 105x
Faster!**

Consistency Levels

STRONG



- Uses **ReadIndex** protocol
- Guarantees linearizability
- Latency: **50-100ms**
- USE CASE: FINANCIAL DATA

LEASE




- Local read if lease valid
- Falls back to Strong if expired
- Latency: **0.5-1ms**
- USE CASE: USER SESSIONS

EVENTUAL



- Always reads local state
- Might return stale data
- Latency: **0.3ms**
- USE CASE: ANALYTICS

LATENCY COMPARISON (LOWER IS BETTER)



LEASE is **109x** faster than STRONG

Why LRU Isn't Enough

STANDARD LRU RULE "Evict the Least Recently Used item"

SCENARIO 1

"holiday-deals"

LRU Assessment

Last access: 2 days ago

EVICT

REALITY CHECK

Will be extremely hot next December (Seasonal Pattern)

SCENARIO 2

"temp-token"

LRU Assessment

Accessed: 1 hour ago

KEEP

REALITY CHECK

One-time use token, never accessed again (Scan Pattern)

THE FUNDAMENTAL FLAW

"LRU only sees **recency**, completely ignoring **access PATTERNS.**"

ML-Based Eviction



CORE IDEA

"Predict which keys will be accessed next"

FEATURE VECTOR

FEATURE NAME	DESCRIPTION
access_count_hour	Number of accesses in last 60 mins
access_count_day	Number of accesses in last 24 hours
total_count	Lifetime access count
time_since_last	Seconds since last access
frequency	Average accesses per hour



RandomForest Classifier



INPUT

Feature vector per cache key (snapshot)



OUTPUT

Probability $P(\text{access})$ in next hour




DECISION

Evict key with **lowest** probability


ML Prediction Algorithm

ml_predict.algo

```
1 INPUT: All cached keys with stats
2 1. Build feature vector for each key
3      # [hour, day, total, time_since, freq]
4 2. Send features to ML model
5 3. Model returns eviction_score
6      # Higher score = More likely to evict
7 4. SELECT key with HIGHEST score
8 5. RETURN key for eviction
```



Eviction Decision Flow



— ML Optimized Path ---- Legacy Fallback

RandomForest Model

WHY RANDOMFOREST

-  Handles **non-linear relationships** in access patterns efficiently
-  Works exceptionally well with **small, tabular datasets**
-  Provides interpretable **confidence scores** for decision thresholds
-  Extremely fast inference (<10µs) suitable for cache operations


CONFIGURATION

PARAMETER	VALUE
n_estimators (Trees)	100
max_depth	10
min_samples_split	5

TRAINING PROCESS


-  TRAINING DATA
Collected from **historical access patterns**
-  LABELS
1 = accessed within next hour, 0 = not accessed
-  LIFECYCLE
Retrains periodically as **workload patterns change**

ML vs LRU Results



METRIC	LRU	ML	CHANGE
Hit Rate Cache efficiency	72.3%	81.7%	+13% ✓
Cold Misses Initial fetch failures	284	183	-36% ✓
Latency Processing overhead	0.5ms	2.1ms	+1.6ms

HIT RATE COMPARISON



Interactive Demo UI

The screenshot shows a dark-themed dashboard with a header bar containing three colored circles (red, yellow, green) and a lock icon followed by the URL "localhost:5173/dashboard".

CLUSTER TOPOLOGY

- Node 1: LEADER
- Node 2: FOLLOWER
- Node 3: FOLLOWER

Below the topology, there are tabs: Election, Cache Data, Performance, and **ML Eviction**. The **ML Eviction** tab is currently active.

LIVE EVICTION CONFIDENCE

KEY	Score	Action
KEY: user:1042	0.98	KEEP
KEY: session:temp	0.12	EVICT
KEY: config:ui	0.85	KEEP

Below the table is a bar chart with six bars. The bars are colored green, green, orange, green, red, and green from left to right. The red bar is significantly shorter than the others.



Real-time Updates

React polling hooks refresh cluster state every 2 seconds



Fault Injection

Kill/Start buttons to simulate node failures instantly



Cache Operations

PUT/GET interface to test replication consistency



Performance Charts

Live latency tracking for Read vs Write operations



ML Visualization

See which keys the model predicts as "hot" or "cold"

Testing Results

Leader Election

✓ PASSED

Follower Recovery

✓ PASSED

Concurrent Writes

✓ PASSED

Read Lease Protocol

✓ PASSED

ML Eviction Accuracy

✓ PASSED



<1 sec

ELECTION TIME



~10 sec

RECOVERY TIME



100%

CONSISTENCY



0%

DATA LOSS

Key Algorithms Summary



LEADER ELECTION

- 01 Follower timeout (Random 150-300ms)
↓
- 02 Become **CANDIDATE** & Request Votes
↓
- 03 Receive Majority Votes
↓
- 04 **Become LEADER**



LOG REPLICATION

- 01 Leader receives Write & Appends to Log
↓
- 02 Send AppendEntries to Followers
↓
- 03 Wait for Majority Acknowledgement
↓
- 04 **COMMIT & Apply to State Machine**



READ LEASE

- 01 Client sends Read Request
↓
- 02 Check: `now < lease_expiry?`

YES (VALID)

Return Local Data
(0.5ms)

NO (EXPIRED)

Run ReadIndex
(50ms)



ML EVICTION

- 01 Cache Full → Collect Key Stats
↓
- 02 RandomForest Model Prediction
↓
- 03 $\text{Score} = P(\text{access_next_hour})$
↓
- 04 **Evict Key with Lowest Score**

Performance Summary

	Election Time	< 1 second
	Failover Downtime	~500ms
	Recovery Time	~10 seconds
	STRONG Read Latency	52ms
	LEASE Read Latency	0.5ms 105x faster
	ML Hit Rate Improvement	+13% vs LRU
	Data Consistency	100%



Strong consistency + High performance

Future Work

Planned roadmap & technical improvements

Log compaction & rotation

Dynamic cluster membership

Multi-region deployment

Prometheus metrics

Security (TLS, authentication)

Production error handling



References

- [1] **In Search of an Understandable Consensus Algorithm**
Diego Ongaro & John Ousterhout, USENIX ATC '14 (Stanford University)
- [2] **Consensus: Bridging Theory and Practice**
Diego Ongaro, PhD Dissertation, Stanford University (2014)
- [3] **The Secret Lives of Data: Raft Visualization**
thesecretlivesofdata.com/raft
- [4] **CoreOS etcd Raft Implementation**
github.com/etcd-io/raft
- [5] **Scikit-learn: Machine Learning in Python**
scikit-learn.org



PROJECT REPOSITORY: github.com/jahnavikedia/raft-cache

Thank You!